

Chapter 10

Association Analysis

Many years ago, a number of new Internet businesses were created to sell books on-line. Over time, they collected information about the books that each of their customers were buying. Using association analysis, they were able to identify groups of books that customers with similar interests seem to have been buying. Using this information, they were able to develop recommendation systems that informed their customers that other customers who purchased some book of interest also purchased other related books.



The customer would often find such recommendations quite useful.

Association analysis identifies relationships or correlations between observations and/or between variables in our datasets. These relationships are then expressed as a collection of so-called association rules. The approach has been particularly successful in mining very large transactional databases, like shopping baskets and on-line customer purchases. Association analysis is one of the core techniques of data mining.

For the on-line bookselling example, historic data is used to identify, for example, that customers who purchased two particular books also tended to purchase another particular book. The historic data might indicate that the first two books are purchased by only 0.5% of all customers. But 70% of these then also purchase the third book. This is an *interesting* group of customers. As a business, we will take advantage

of this observation by targeting advertising of the third book to those customers who have purchased both of the other books.

The usual type of input data for association analysis consists of transactional data, such as the items in a shopping basket at a supermarket, books and videos purchased by a single client, or medical treatments and tests received by patients. We are interested in the items whose co-occurrence within a transaction is of interest.

This short chapter introduces the basic concepts of association analysis and how to perform it in both Rattle and R.

10.1 Knowledge Representation

A representation of association rules is required to identify relationships between items within transactions. Suppose each transaction is thought of as a basket of items (which we might represent as $\{A, B, C, D, E, F\}$). The aim is to identify collections of items that appear together in multiple baskets (e.g., perhaps the items $\{A, C, F\}$ appear together in quite a few shopping baskets). From these so called *itemsets* (i.e., sets of items) we identify rules like $A, F \Rightarrow C$ that tell us that when A and F appear in a transaction (e.g., a shopping basket) then typically so does C .

A collection of association rules then represents a model as the outcome of association analysis. The general format of an association rule is

$$\mathcal{A} \rightarrow \mathcal{C}.$$

Both \mathcal{A} (the left hand side or antecedent) and \mathcal{C} (the right side or consequent) are sets of items. Generally we think of items as being particular books, for example, or particular grocery items. Examples might be:

$$milk \rightarrow bread,$$

$$beer \ \& \ nuts \rightarrow potato \ crisps,$$

$$Shrek1 \rightarrow Shrek2 \ \& \ Shrek3.$$

The concept of an item can be generalised to a specific variable/value combination as the item. The concept of association analysis can then be applied to many different datasets. Using our *weather* dataset, for example, this representation will lead to association rules like

$$WindDir3pm = NNW \rightarrow RainToday = No.$$

10.2 Search Heuristic

The basis of an association analysis algorithm is the generation of frequent itemsets. A frequent itemset is a set of items that occur together frequently enough to be considered as a candidate for generating association rules.

The obvious approaches to identifying itemsets that appear frequently enough in the data are quite expensive computationally, even with moderately sized datasets. The apriori algorithm takes advantage of the simple observation that all subsets of a frequent itemset must also be frequent. That is, if $\{milk, bread, cheese\}$ is a frequent itemset then so must each of the smaller itemsets, $\{milk, bread\}$, $\{milk, cheese\}$, $\{bread, cheese\}$, $\{milk\}$, $\{bread\}$, and $\{cheese\}$. This observation allows the algorithm to consider a significantly reduced search space by starting with frequent individual items. This first step eliminates very rare items. We then combine the remaining single items into itemsets containing just two items and retain only those that are frequent enough and similarly for itemsets containing three items and so on.

The concept of *frequent enough* is a parameter of the algorithm used to control the number of association rules discovered. This is called the *support* and specifies how frequently the items must appear in the whole dataset before they can be considered as a candidate association rule. For example, the user may choose to consider only sets of items that occur in at least 5% of all transactions.

The second phase of the algorithm considers each of the frequent itemsets and for each generates all possible combinations of association rules. Thus, for an itemset containing three items $\{milk, bread, cheese\}$, the following are among the possible association rules that will be considered:

$$\begin{aligned} bread \ \& \ milk &\rightarrow cheese, \\ milk &\rightarrow bread \ \& \ cheese, \\ cheese \ \& \ milk &\rightarrow bread, \end{aligned}$$

and so on.

The actual association rules that we retain are those that meet a criterion called *confidence*. The confidence calculates the proportion of transactions containing \mathcal{A} that also contain \mathcal{C} . The confidence specifies a minimal probability for the association rule. For example, the user may choose to generate only rules that are true at least 90% of the time (that

is, when \mathcal{A} appears in the basket, \mathcal{C} also appears in the same basket at least 90% of the time).

The apriori algorithm is a breadth-first or generate-and-test type of search algorithm. Only after exploring all of the possibilities of associations containing k items does it then consider those containing $k + 1$ items. For each k , all candidates are tested to determine whether they have enough support.

In summary, the algorithm uses a simple two-phase generate-and-merge process. In phase 1, we generate frequent itemsets of size k (iterating from 1 until we have no frequent k -itemsets) and then combine them to generate candidate frequent itemsets of size $k + 1$. In phase 2, we build candidate association rules.

10.3 Measures

The two primary measures used in association analysis are the **support** and the **confidence**. The minimum support is expressed as a percentage of the total number of transactions in the dataset. Informally, it is simply how often the items appear together from amongst all of the transactions. Formally, we define *support* for a collection of items \mathcal{I} as the proportion of all transactions in which all items in \mathcal{I} appear and express the support for an association rule as

$$\text{support}(\mathcal{A} \rightarrow \mathcal{C}) = P(\mathcal{A} \cup \mathcal{C}).$$

Typically, we use small values for the support, since overall the items that appear together “frequently” enough that are of interest generally won’t be the obvious ones that regularly appear together.

The minimum confidence is also expressed as the proportion of the total number of transactions in the dataset. Informally, it is a measure of how often the items \mathcal{C} appear whenever the items \mathcal{A} appear in a transaction. Formally, it is a conditional probability:

$$\text{confidence}(\mathcal{A} \rightarrow \mathcal{C}) = P(\mathcal{C}|\mathcal{A}) = P(\mathcal{A} \cup \mathcal{C})/P(\mathcal{A}).$$

It can also be expressed in terms of the *support*:

$$\text{confidence}(\mathcal{A} \rightarrow \mathcal{C}) = \text{support}(\mathcal{A} \rightarrow \mathcal{C})/\text{support}(\mathcal{A}).$$

Typically, this measure will have larger values since we are looking for the association rules that are quite strong, so that if we find the items in \mathcal{A} in a transaction then there is quite a good chance of also finding \mathcal{C} in the transaction.

There are a collection of other measures that are used with association rule analysis. One that is used in R and hence Rattle is the **lift**. The lift is the increased likelihood of \mathcal{C} being in a transaction if \mathcal{A} is included in the transaction. It is calculated as

$$lift(\mathcal{A} \rightarrow \mathcal{C}) = confidence(\mathcal{A} \rightarrow \mathcal{C}) / support(\mathcal{C}).$$

Another measure is the **leverage**, which captures the fact that a higher frequency of \mathcal{A} and \mathcal{C} with a lower lift may be interesting:

$$leverage(\mathcal{A} \rightarrow \mathcal{C}) = support(\mathcal{A} \rightarrow \mathcal{C}) - support(\mathcal{A}) * support(\mathcal{C}).$$

10.4 Tutorial Example

Two types of association rules were identified above, corresponding to the type of data made available. The simplest case, known as market basket analysis, is when we have a transaction dataset that records just a transaction identifier. The identifier might identify a single shopping basket containing multiple items from shopping or a particular customer or patient and their associated purchases or medical treatments over time. A simple example of a market basket dataset might record the purchases of DVDs by customers (three customers in this case):

```
ID,Item
1,Sixth Sense
1,LOTR1
1,Harry Potter1
1,Green Mile
1,LOTR2
2,Gladiator
2,Patriot
2,Braveheart
3,LOTR1
3,LOTR2
```

The resulting model will then be a collection of association rules that might include

$$LOTR1 \rightarrow LOTR2.$$

The second form of association rule uses a dataset that we are more familiar with. This approach treats each observation as a transaction and the variables as the items in the “shopping basket.” Considering the *weather* dataset, we might obtain models that include rules of the form

$$Humidity3pm = High \ \& \ Pressure3pm = Low \rightarrow RainToday = Yes.$$

Both forms are supported in Rattle and R.

Building a Model Using Rattle

Rattle builds association rule models through the **Associate** tab. The two types of association rules are supported and the appropriate type is chosen using the **Baskets** check button. If the button is checked then Rattle will use the **Ident** and **Target** variables for the analysis, performing a market basket analysis. If the button is not checked, then Rattle will use the **Input** variables for a rules analysis.

For a basket analysis, the data is thought of as representing shopping baskets (or any other type of collection of items, such as a basket of medical tests, a basket of medicines prescribed to a patient, a basket of stocks held by an investor, and so on). Each basket has a unique identifier, and the variable specified as an **Ident** variable on the **Data** tab is taken as the identifier of a shopping basket. The contents of the basket are then the items contained in the column of data identified as the **Target** variable. For market basket analysis, these are the only two variables used.

To illustrate market basket analysis with Rattle, we can use a very simple and trivial dataset consisting of the DVD movies purchased by customers. The data is available as a CSV file (named `dvdtrans.csv`) from the Rattle package. The simplest way to load this dataset into Rattle is to first load the default sample *weather* dataset from the `weather.csv` file into Rattle. We do this by clicking the **Execute** button on starting Rattle. Then click the **Filename** button (which will now be showing `weather.csv`) to list the contents of Rattle’s sample CSV folder. Choose `dvdtrans.csv` and click **Open** and then **Execute**. The ID variable will

automatically be chosen as the `Ident`, but we will need to change the role of `Item` to be `Target`, as in Figure 10.1.




Figure 10.1: Choose the `dvdtrans.csv` file and load it into Rattle with a click of the `Execute` button. Then set the role for `Item` to be `Target` and click `Execute` for the new role to be noted.

On the `Associate` tab, ensure that the `Baskets` button is checked. Click the `Execute` button to build a model that will consist of a collection of association rules. Figure 10.2 shows the resulting text view, which we now review.

The first few lines of the text view list the number of association rules that make up the model. In our example, there are 127 rules:

```
Summary of the Apriori Association Rules:

Number of Rules: 127
```

The next code block reports on the distribution of the three measures as found for the 127 rules of the model:

```
Summary of the Measures of Interestingness:

      support      confidence      lift
Min.   :0.100   Min.   :0.100   Min.   : 0.714
1st Qu.:0.100   1st Qu.:0.500   1st Qu.: 1.429
Median :0.100   Median :1.000   Median : 2.500
Mean   :0.145   Mean   :0.759   Mean   : 3.015
3rd Qu.:0.100   3rd Qu.:1.000   3rd Qu.: 5.000
Max.   :0.700   Max.   :1.000   Max.   :10.000
```

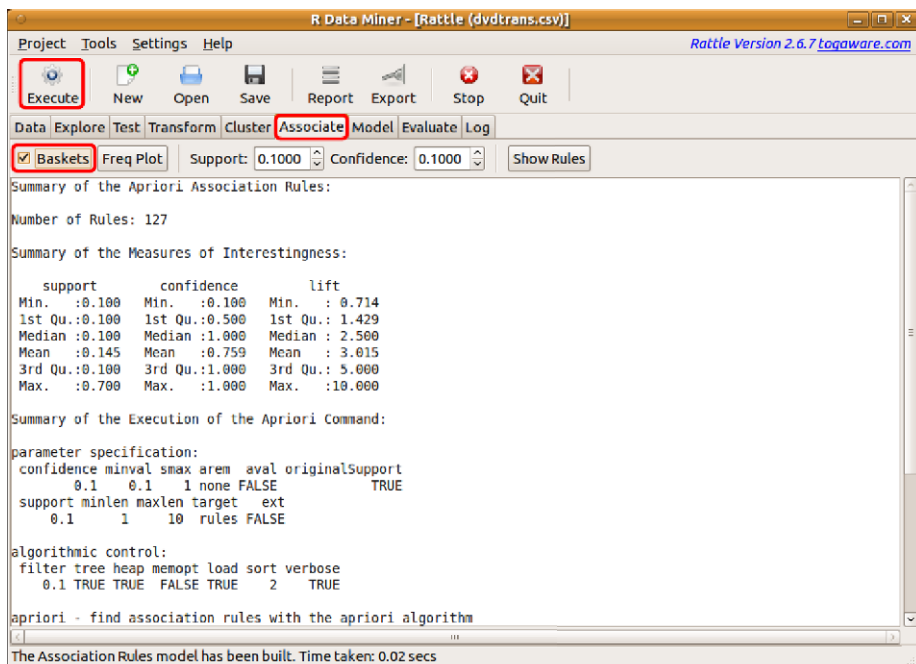


Figure 10.2: Building an association rules model.

The 127 association rules met the criteria of having a minimum support of 0.1 and a minimum confidence of 0.1. Across the rules the support ranges from 0.1 up to 0.4. Confidence ranges from 0.1 up to 1.0 and lift from 0.83 up to 10.0.

This section is followed by a summary of the process of building the model. It begins with a review of the options supplied or the default values for the various parameters. We can see `confidence=` and `support=` listed:

Summary of the Execution of the Apriori Command:

parameter specification:

```
confidence minval smax arem  aval originalSupport
      0.1      0.1      1 none FALSE              TRUE
support minlen maxlen target  ext
      0.1       1     10  rules FALSE
```


These options are tunable through the Rattle interface. Others can be tuned directly through R. A set of parameters that control how the algorithm itself operates is then displayed:

```
algorithmic control:
  filter tree heap memopt load sort verbose
    0.1 TRUE TRUE FALSE TRUE    2    TRUE
```

The final section includes detailed information about the algorithm and the model that has been built:

```
apriori - find association rules with the apriori algorithm
version 4.21 (2004.05.09) (c) 1996-2004 Christian Borgelt
set item appearances ... [0 item(s)] done [0.00s].
set transactions ... [10 item(s), 10 trans] done [0.00s].
sorting and recoding items ... [10 item(s)] done [0.00s].
creating transaction tree ... done [0.00s].
checking subsets of size 1 2 3 4 5 done [0.00s].
writing ... [127 rule(s)] done [0.00s].
creating S4 object ... done [0.00s].
```

The Show Rules button will show all of the association rules for the model in the text view window, sorted by the level of confidence in the rule. The top five rules will be:

lhs	rhs	supp	conf	lift
1 {Harry Potter2}	=> {Harry Potter1}	0.1	1	5.000
2 {Braveheart}	=> {Patriot}	0.1	1	1.667
3 {Braveheart}	=> {Gladiator}	0.1	1	1.429
4 {LOTR}	=> {Green Mile}	0.1	1	5.000
5 {LOTR}	=> {Sixth Sense}	0.1	1	1.667

These rules have only a single item on each side of the arrow, and all have a support of 0.1 and a confidence of 1. We can see that for either of the first two movies there is quite a large lift obtained.

Building a Model Using R

Arules (Hahsler et al., 2011) provides `apriori()` for R. The package provides an interface to the widely used, and freely available, apriori

software from Christian Borgelt. This software was, for example, commercially licensed for use in the Clementine¹ data mining package and is a well-developed and respected implementation.

When loading a dataset to process with `apriori()`, it needs to be converted into a transaction data structure. Consider a dataset with two columns, one being the identifier of the “basket” and the other being an item contained in the basket, as is the case for the `dvdtrans.csv` data. We can load that data into R:

```
> library(arules)
> library(rattle)
> dvdtrans <- read.csv(system.file("csv", "dvdtrans.csv",
                                   package="rattle"))
> dvdDS <- new.env()
> dvdDS$data <- as(split(dvdtrans$Item, dvdtrans$ID),
                   "transactions")
> dvdDS$data
transactions in sparse format with
 10 transactions (rows) and
 10 items (columns)
```

We can then build the model using this transformed dataset:

```
> dvdAPRIORI <- new.env(parent=dvdDS)
> evalq({
  model <- apriori(data, parameter=list(support=0.2,
                                         confidence=0.1))
}, dvdAPRIORI)
```

The rules can be extracted and ordered by confidence using `inspect()`. In the following code block we also use `[1:5]` to limit the display to just the first five association rules. We notice that the first two are symmetric, which is expected since everyone who purchases one of these movies always also purchases the other.

¹Clementine became an SPSS product and was then purchased by IBM to become IBM SPSS Modeler.

```
> inspect(sort(dvdAPRIORI$model, by="confidence")[1:5])
```

	<i>lhs</i>	<i>rhs</i>	<i>support</i>	<i>confidence</i>	<i>lift</i>
1	{ <i>LOTR1</i> }	=> { <i>LOTR2</i> }	0.2	1	5.000
2	{ <i>LOTR2</i> }	=> { <i>LOTR1</i> }	0.2	1	5.000
3	{ <i>Green Mile</i> }	=> { <i>Sixth Sense</i> }	0.2	1	1.667
4	{ <i>Patriot</i> }	=> { <i>Gladiator</i> }	0.6	1	1.429
5	{ <i>Patriot</i> , <i>Sixth Sense</i> }	=> { <i>Gladiator</i> }	0.4	1	1.429

10.5 Command Summary

This chapter has referenced the following R packages, commands, functions, and datasets:

<code>apriori()</code>	function	Build an association rule model.
<code>arules</code>	package	Support for association rules.
<code>inspect()</code>	function	Display results of model building.
<code>weather</code>	dataset	Sample dataset from rattle .